

Exercise 0 Dataset description

e12045110 Maria de Ronde
e12045110 Quentin Andre
e11921655 Fabian Holzberger

March 23, 2021

Classification Dataset: Email-Spam ([link to dataset](#))

We have chosen a e-mail-spam dataset for classification. The link above contains three datasets from which choose two, namely the `lingSpam.csv` and `enormSpamSubset.csv` for our project, since they have no missing values and the same layout. The goal is to distinguish by a machine learning algorithm between spam and non-spam e-mails. The structure of the datasets is shown in table 1. Here we see that the datasets have a Body-column, that contains the text-body of an e-mail and a Label-column, that is either set to 1 for spam e-mails or 0 for non-spam e-mails. Since the text-data is not structured, we can't assign it to a level of measurement. The Label on the other side is nominal. For the rest of this document we assume the two datasets are concatenated

Index	Body	Label
100	Subject: inexpensive online medication here pummel wah springtail cutler bodyguard we ship quality medications overnight to your door !...	1
6006	Subject: organizational changes we are pleased to announce the following organizational changes : enron global assets and services in order to increase senior management focus on our international businesses...	0

Figure 1: Structure of the Email-Spam Dataset

into one big dataset. Note that the last row of the `lingSpam.csv` contains a summary of all its previous data that we removed. The dataset then contains in total 12604 samples, of which 327 are duplicates, where 42% of the samples are spam- and 58% are non-spam e-mails. The duplicates are removed in the analysis of this report. The length of emails in the dataset is shown in figure 2 a). The length of e-mails ranges is between 10 and 121502 characters wherein most have a length of 1000 characters.




Figure 2: a) Distribution of e-mail lengths b) 10 most common words in spam e-mails c) 10 most common words in non-spam e-mails

We aim to apply the *Bag of Words* method to the dataset. This method extracts the N most common words from all e-mails and then maps an e-mail to a vector v , such that the component m of v is a non-negative integer, that counts the occurrences of the m th most common word in the corresponding e-mail. From that we conclude that the dimension of our dataset is $N + 1$ after the bag of words transformation, since we also include the target attribute. Note that we apply the following cleanup steps to all e-mails to remove data, that we expect to not improve the classification: 1. remove links, 2. remove characters except alphabetical ones, 3. convert uppercase-chars into lowercase-chars, 4. lemmatize words 5. remove stopwords. By that we reduce

the number of total distinct words in all e-mails from 126019 words to 103759 words (see Jupyter notebook). This also illustrates that our classification algorithm has to work with data that is high dimensional (dimension much larger than 1000 to represent e-mails sufficiently). Since the raw data is unstructured we can't assign it to a level of measurement or dimension. In figure 2 b) and c) we show the most common words in spam and non-spam emails after applying our cleanup steps. One can see that the word count in spam e-mails is much higher than in the non-spam e-mails, which means that overall spam mails have more words in common than non-spam e-mails.

Regression Dataset: Automobile ([link to dataset](#))

The automobile dataset is used for regression. It consists of 205 samples and 26 attributes where 16 attributes are given as numerical values and 10 further as nominal attributes in text form. In figure 3 a) we show the numerical attributes in tabular form. In the figure we see that the symbolizing attribute is ordinal, where the others are all ratios. In the count column we note that some attributes are given for less than 205 cars, meaning that this data are missing, see also the last plot of figure 4 for a distribution of missing values. For four rows the target value is missing, these rows are excluded from the data. Normalized losses has 41 missing values and will be excluded from the model. For the rest of the missing values we intend to apply a imputation technique. This is a crucial step, since the dataset is small and deleting cars with incomplete attributes would lead to significant information loss.

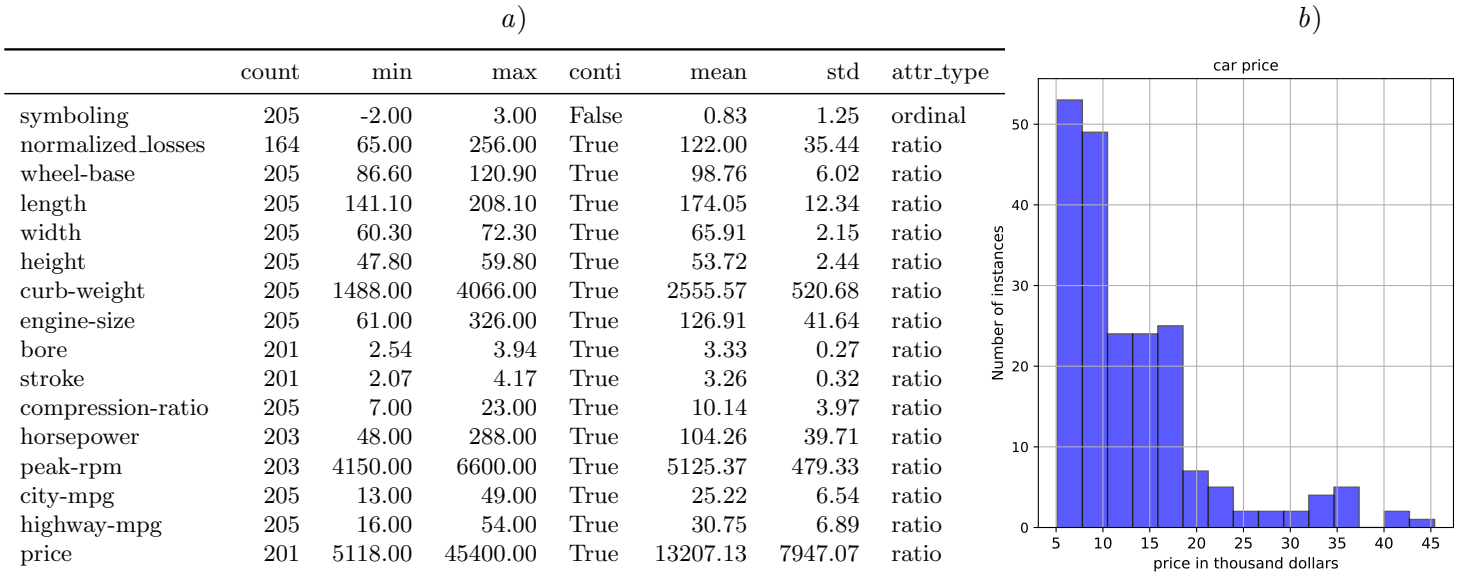


Figure 3: Numerical attributes: a) numerical attributes of the automobile dataset b) distribution of target variable

The target value for the regression is the price of a car, which is a continuous variable between 5118 and 45400 dollars as can be seen also in 3 a). In figure 3 b) the distribution of the price for a car is shown. Here we note that the price is not evenly distributed over its range and that more than 80% of the cars cost between 5.000 and 20.000 dollar. Cars more expensive than 45.000 will most likely not be predicted properly, because they are not represented in the dataset. By comparing the ranges of different numerical attributes, we note that some of them differ in orders of magnitude like the stroke and curb-weight, that differ by about two magnitudes. Therefore, we might re-scale them (for example normalizing) to improve the performance of our machine learning algorithm. Note that re-scaling of the target attribute is not necessary for the most

algorithms. Additionally many attributes are tail heavy, meaning their mean value differs significantly from their max and minimum sample-value. We can also try to reshape attributes that are tail-heavy into more evenly shaped distributions by a transformation to improve the performance of our algorithms. The nominal attributes are plotted as histograms in figure 4. Since the nominal attributes are categorical without a defined order, we will encode them to N binary attributes.

Comparison of the datasets

In table 5 we compare the two datasets from above.

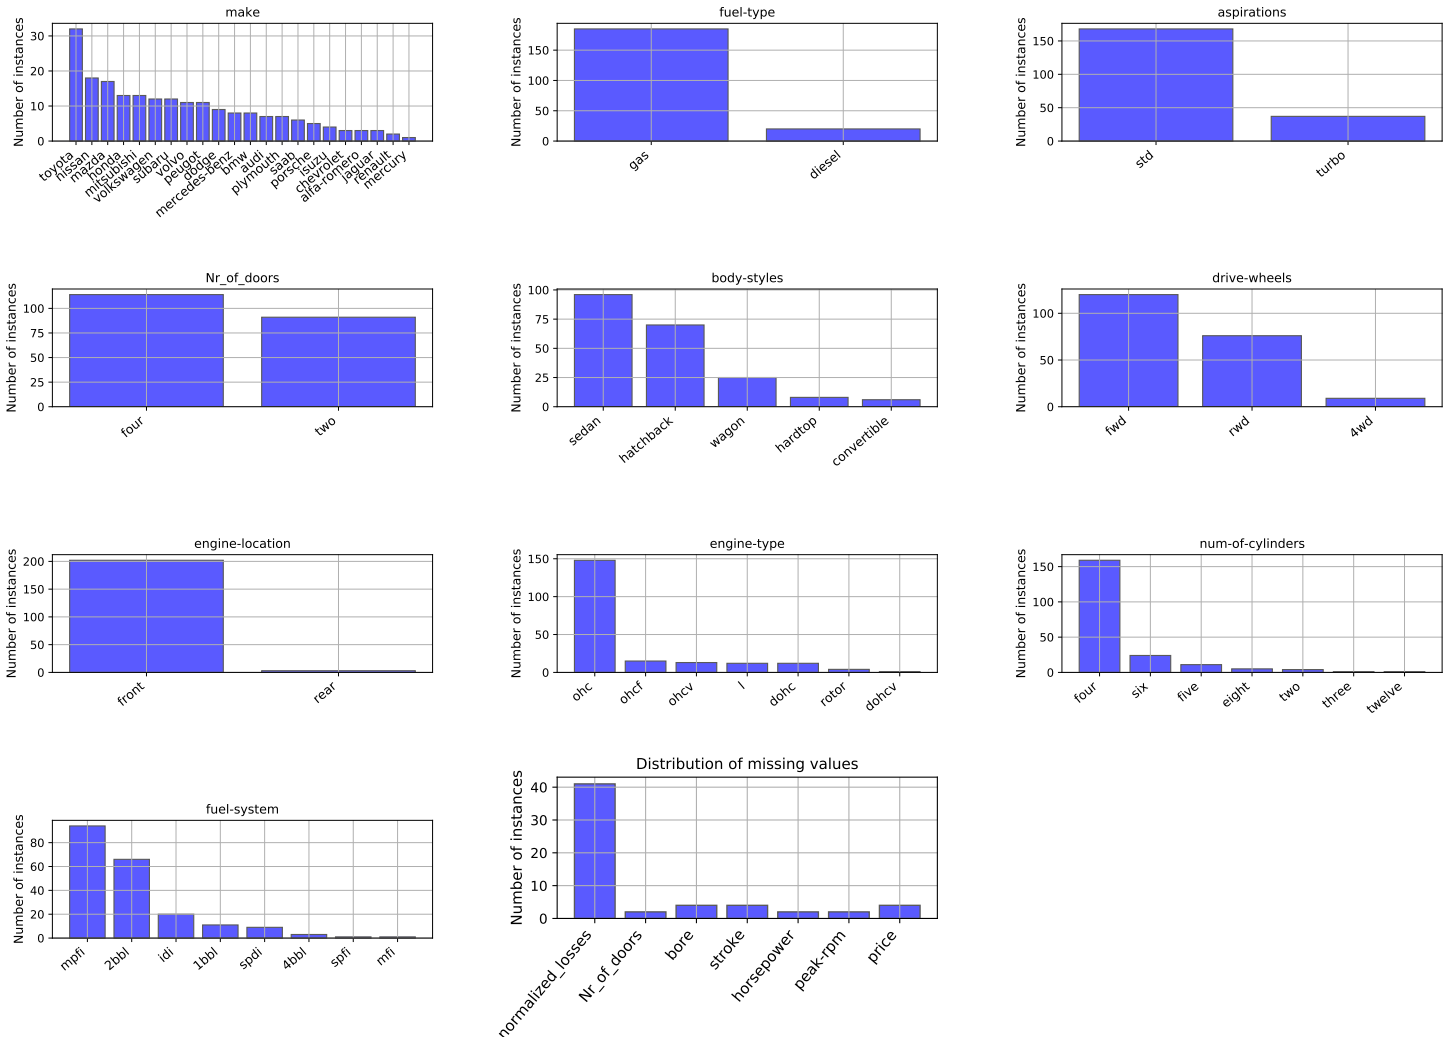


Figure 4: Nominal attributes of the automobile dataset where the last graph summarizes the missing values over all attributes

Compared to the automobile dataset, the e-mail-spam dataset has about 60 times more samples, such that we consider it as large dataset, where the automobile dataset is a small dataset. When considering the attribute-types one notes, that the automobile dataset contains all levels of measurement, where the e-mail spam dataset contains only nominal attributes. Further since the latter one is not a structured dataset it does not have a dimension, compared to the dimension of 26 in the automobile dataset. In the automobile dataset some

Dataset	samples	attribute-types	dimension	data-types	missing Values	duplicates
E-mail Spam	12.277	nominal	-	text, boolean	no	yes
Automobile	205	nominal, ordinal, intervall, ratio	26	integer, real, text	yes	no

Figure 5: Characteristics of the chosen datasets

attributes are missing, compared to the other one where no values are missing but data is duplicated. By this we conclude that the two datasets differ greatly, as required for this exercise.