

Exercise0 Dataset description

e12045110 Maria de Ronde
e Quentin Andre
e11921655 Fabian Holzberger

March 17, 2021

Classification Dataset: Email-Spam

We have chosen a e-mail-spam dataset for classification ([link to dataset](#)). The goal for this dataset is to distinguish by a machine learning algorithm between spam and non-spam e-mails. The dataset is given in `csv` format and the structure is shown in table 1. Here we see that the dataset has a nominal Body-column that contains the text-body of an email and a nominal Label-column that is either set to 1 for spam e-mails or 0 for non-spam e-mails. The dataset contains in total 12278 samples, of which 327 duplicates, where 42%

Index	Body	Label
100	Subject: inexpensive online medication here pummel wah springtail cutler bodyguard we ship quality medications overnight to your door !...	1
6006	Subject: organizational changes we are pleased to announce the following organizational changes : enron global assets and services in order to increase senior management focus on our international businesses...	0

Figure 1: Structure of the Email-Spam Dataset

of the samples are spam- and 58% are non-spam e-mails. We aim to apply the **Bag of Words** method to the dataset. This method extracts the N most common words from all E-mails and then maps an e-mail to a vector v , such that the component m of v is an non-native integer, that counts the occurrences of the m th most common word in the corresponding e-mail. From that we conclude that the dimension of our dataset is $N + 1$ after the Bag of words transformation, since we also include the target attribute. Note that we apply the following cleanup steps to all e-mails to remove data, that we expect to not improve the classification: 1. remove links, 2. remove characters except alphabetical ones, 3. convert uppercase-chars into lowercase-chars, 4. lemmatize words 5. remove stopwords. By that we reduce the number of total distinct words in all e-mails from 126704 words to 103784 words (see jupyter notebook). This also illustrates that our classification algorithm has to work with data that is high dimensional (dimension much larger than 1000). Since the raw data merely contain the e-mail bodies and spam/non-spam attributes, we assume in this document a dimension of two for this data.




Figure 2: Figure a) 15 most common words in spam e-mails b) 15 most common words in non-spam e-mails

In figure 2 we show the most common words in spam and non-spam emails. One can see that the word count in spam e-mails is much higher than in the non-spam e-mails which means that overall spam mails have more words in common than non-spam e-mails.

Regression Dataset: Automobile

The automobile dataset is used for regression ([link to dataset](#)). It consists of 205 samples and 26 attributes where 16 attributes are given as numerical values and 10 further as nominal attributes in text form. In figure 3 a) we show the numerical attributes in tabular form. In the figure we see that the symbolizing attribute is ordinal, where the others are all ratios. In the count column we note that some attributes are given for less than 205 cars, meaning that this data are missing, see also the last plot of figure 4 for a distribution of missing values. For four rows the target value is missing, these rows will be excluded from the data. Normalized losses has 41 missing values and will be excluded from the model. For the rest of the missing values we intend to apply a imputation technique. This is a crucial step, since the dataset is small and deleting cars with incomplete attributes would lead to significant information loss.

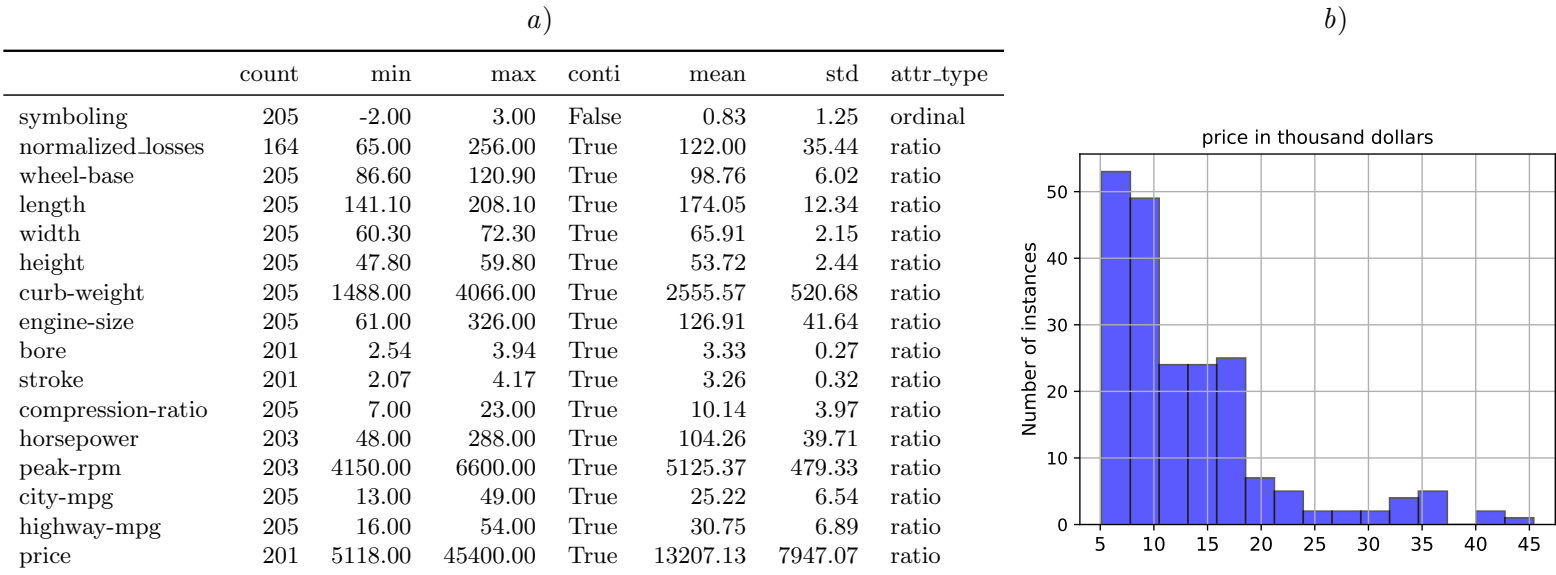


Figure 3: Numerical attributes: a) numerical attributes of the automobile dataset b) distribution of target variable

The target value for the regression is the price of a car, which is a continuous variable between 5118 and 45400 dollars as can be seen also in 3 a). In figure 3 b) the distribution of the price for a car is shown. Here we note that the price is not evenly distributed over its range and that more than 80% of the cars cost between 5.000 and 20.000 dollar. Cars more expensive than 45.000 will most likely not be predicted properly, because they are not represented in the dataset. By comparing the ranges of different numerical attributes, we note that some of them differ in orders of magnitude like the stroke and curb-weight, that differ by about two magnitudes. Therefore, we might rescale them (for example normalizing) to improve the performance of our machine learning algorithm. Note that re-scaling of the target attribute is not necessary for the most algorithms. Additionally many attributes are tail heavy, meaning their mean value differs significantly from their max and minimum sample-value. We can also try to reshape attributes that are tail-heavy into more evenly shaped distributions by a transformation to improve the performance of our algorithms. The nominal attributes are plotted as histograms in figure 4. Since the nominal attributes are categorical without a defined order, we will encode them to N binary attributes.

Comparison of the datasets

In table 5 we compare the two datasets from above.

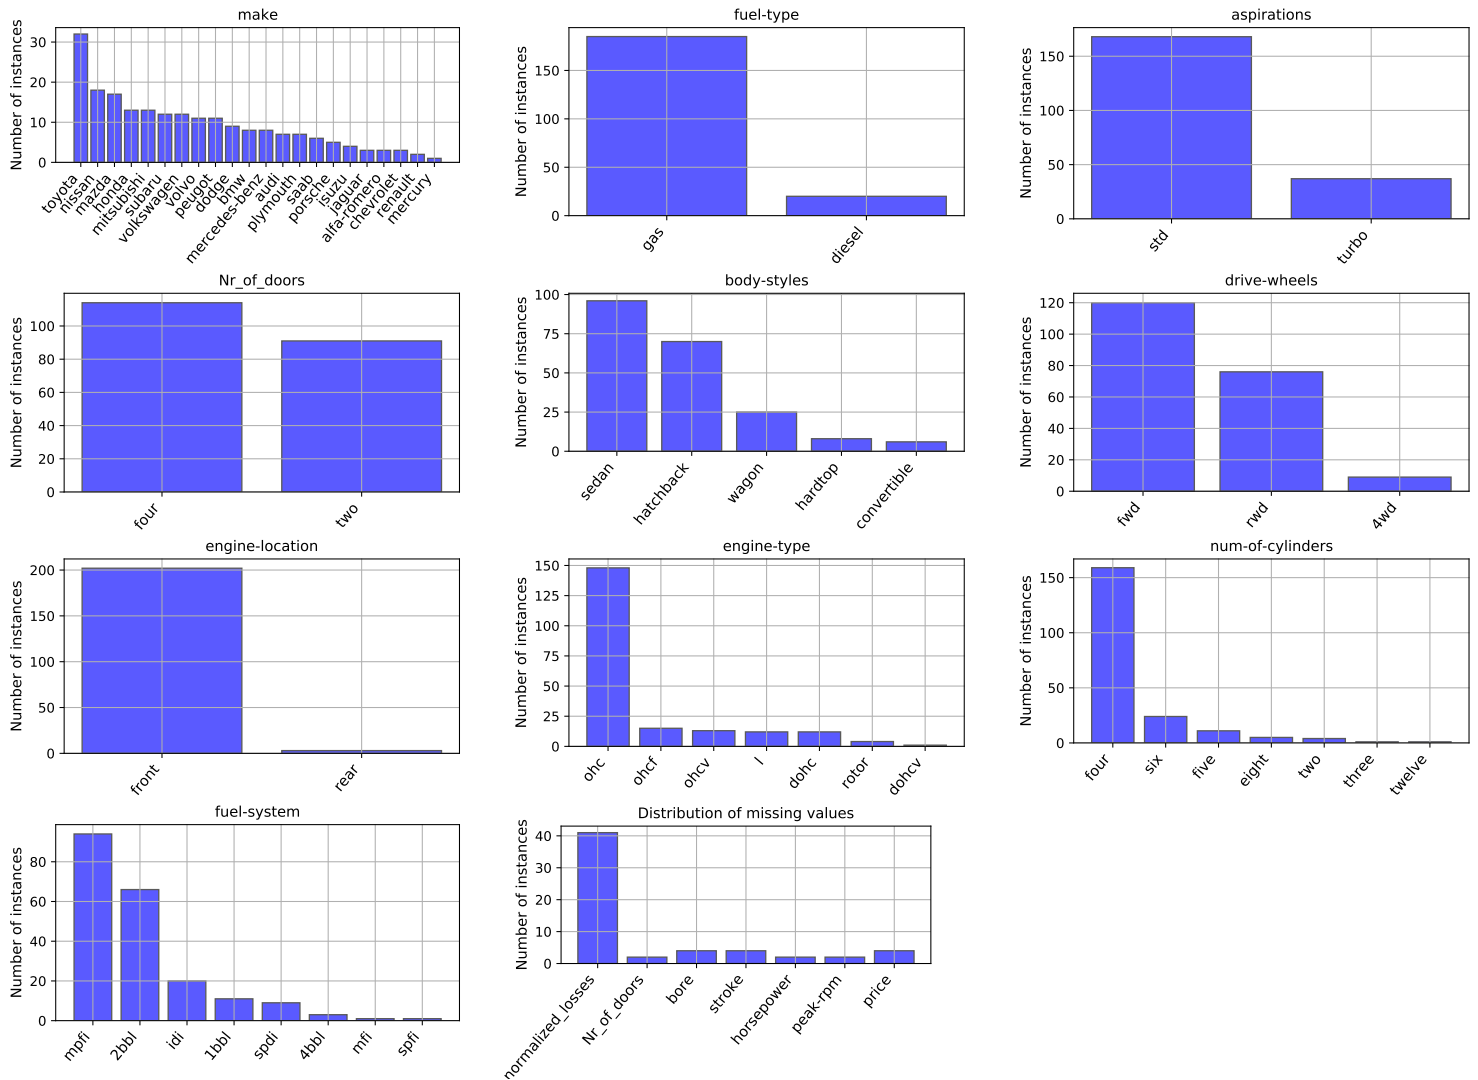


Figure 4: Nominal attributes of the automobile dataset where the last graph summarizes the missing values over all attributes

Dataset	samples	attribute-types	dimension	attribute format	missing Values
E-mail Spam	10.000	nominal	$N + 1$	text, boolean	no
Automobile	205	nominal, ordinal,intervall,ratio	26	integer, real, text	yes

Figure 5: Characteristics of the chosen datasets

Compared to the Automobile dataset, the e-mail-spam dataset has about 50 times more samples such that we consider it as large dataset, where the automobile dataset is a small dataset. When considering the attribute-types one notes, that the automobile dataset contains all levels of measurement, where the e-mail spam dataset contains only nominal attributes. Further the dimension of the later one is two as stated before, which is rather low compared to the dimension of 26 in the Automobile dataset. In the Automobile

dataset some attributes are missing, compared to the other one where no values are missing. By this we conclude that the two datasets differ greatly, as required for this exercise.